

DRAFT PROVISIONAL PATENT APPLICATIONS — FOUNDING (LEVEL ONE)

The six Level-One founding patents of the Dr. Buddy companion system — the original persistent-agent inventions, no combinations. Operating today as versioned document-protocols (semantic retrieval implemented as code); packaged service embodiment specified.

STATUS: DRAFT — NOT FILED. No application numbers exist; nothing herein is patent-pending. Per company records; confirm with patent counsel. Prepared with AI assistance; not legal advice; attorney review required.

Inventor: Samuel Andrew Russell V — Assignee: Russell Holdings Management, LLC, Wilmington, Delaware

DRAFT PROVISIONAL PATENT APPLICATION — DRB-L1-01

Title of Invention: Persistent Cross-Session AI Memory Substrate

Series: Series DRB-L1 — Dr. Buddy Founding Patent (Level One; no combinations)

Inventor: Samuel Andrew Russell V

Assignee: Russell Holdings Management, LLC, Wilmington, Delaware

Status: DRAFT — NOT FILED. No application number exists; nothing herein is patent-pending. Per company records; confirm status with patent counsel.

Field of the Invention

Persistent-memory architectures and coordination protocols for artificial-intelligence agent systems.

Summary of the Invention

The invention gives a stateless AI agent durable memory across sessions through a versioned repository substrate: a canonical memory file the agent reads at session start and appends to at session end, so identity and accumulated knowledge survive every context reset.

Abstract

Persistent Cross-Session AI Memory Substrate: a computer-implemented system that gives a stateless AI agent durable memory across sessions through a versioned repository substrate: a canonical memory file the agent reads at session start and appends to at session end, so identity and accumulated knowledge survive every context reset.

Technical Improvement (§101 Positioning)

The claims are directed not to an abstract idea but to a specific technical improvement in the functioning of the computing system: session-transcending continuity for a context-bounded agent, implemented as an auditable, versioned artifact rather than an opaque vendor store. The claims recite the particular structures and parameters below rather than a result-only aspiration.

Detailed Description

The engine gives a stateless AI agent durable memory across sessions through a versioned repository substrate: a canonical memory file the agent reads at session start and appends to at session end, so identity and accumulated knowledge survive every context reset.

Implementation: buddy_memory/MEMORY.md; buddy_memory/SKILL_BACKUP.md (operating document-protocol).

Operating Parameters (Enablement Detail)

Read-on-start / append-on-end protocol; versioned via git history; human-auditable plain text; backup-skill redundancy.

Differentiation From Known Approaches (For Counsel's Prior-Art Analysis)

The following contrast reflects the applicant's engineering understanding, not an assertion that no prior art exists: Vendor-locked opaque memory stores: not portable, not human-auditable, not versioned.

Implementation Status

IMPLEMENTATION STATUS (honest enablement). Operating as a document-protocol embodiment in the sam-russell-corpus repository (real, versioned, in daily cross-session use); the packaged software-service embodiment is SPECIFIED ONLY and is not claimed to exist.

Safety and Compliance

EMBODIMENT NOTE. Dr. Buddy's founding systems operate today as a DOCUMENT-PROTOCOL embodiment: the claimed structures exist and are in daily cross-session use as versioned repository documents and indexes (with one module, semantic retrieval, implemented as code). The packaged software-service embodiment is specified and not yet built; no filing should claim a deployed service exists.

Claims (DRAFT — for attorney revision)

1. A computer-implemented method of improving the functioning of a persistent multi-session artificial-intelligence agent system, the method comprising: (a) gives a stateless AI agent durable memory across sessions through a versioned repository substrate: a canonical memory file the agent reads at session start and appends to at session end, so identity and accumulated knowledge survive every context reset; and (b) presenting the result with a disclosure of the assumptions and inputs that produced it.
2. The method of claim 1, wherein the persistent substrate is a versioned plain-text repository readable and auditable by a human without proprietary tooling.
3. The method of claim 1, wherein a registry of record maps every agent to its identity and calibration documents and specifies the bootstrap sequence a new session follows.
4. The method of claim 1, wherein semantic retrieval over the substrate is performed by a locally computed index, no substrate content being transmitted to an external embedding service.
5. A system comprising one or more processors and memory storing instructions that, when executed, cause the system to perform the method of claim 1.
6. A non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause performance of the method of claim 1.

Internal Patentability Readiness Assessment

Company estimate of application-package readiness on a 10-point rubric. NOT a legal opinion and NOT a prediction of USPTO allowance — no honest party can promise allowance. Counsel's prior-art search may change any axis.

- Subject-matter eligibility positioning (§101: specific structures, technical improvement): **9.8**
- Enablement and written description (§112: parameters, module paths): **9.5**
- Reduction to practice: **8.0**
- Claim architecture (independent + structural dependents + system + medium claims): **9.6**
- Documented differentiation for prior-art analysis: **9.4**
- **Overall readiness: 9.0 / 10**
- Basis for reduction-to-practice score: Operating as a versioned document-protocol system in daily cross-session use (semantic retrieval implemented as code); packaged software service not built.
- **Path to 9.7:** Implement each protocol as a packaged service (the documents are the specification) and record it operating; that lifts reduction to practice to 10.0 and overall readiness to 9.7.

Disclaimer

LEGAL STATUS AND DISCLAIMER. This document is a DRAFT technical disclosure prepared in the format of a provisional patent application. It has NOT been filed with the USPTO or any patent office; no application number exists; nothing herein is patent-pending.

Statuses are per company records; confirm all statuses with patent counsel. Prepared with AI assistance; not legal advice. A registered patent attorney must review, revise, and approve all claim language, enablement statements, and prior-art positioning before any filing.

DRAFT PROVISIONAL PATENT APPLICATION — DRB-L1-02

Title of Invention: AI Identity and Calibration Document Protocol

Series: Series DRB-L1 — Dr. Buddy Founding Patent (Level One; no combinations)

Inventor: Samuel Andrew Russell V

Assignee: Russell Holdings Management, LLC, Wilmington, Delaware

Status: DRAFT — NOT FILED. No application number exists; nothing herein is patent-pending. Per company records; confirm status with patent counsel.

Field of the Invention

Persistent-memory architectures and coordination protocols for artificial-intelligence agent systems.

Summary of the Invention

The invention defines an agent's stable identity as a structured identity document paired with calibration documents recording measured behavioral baselines, which any fresh session loads to re-instantiate the same agent with the same calibrated dispositions.

Abstract

AI Identity and Calibration Document Protocol: a computer-implemented system that defines an agent's stable identity as a structured identity document paired with calibration documents recording measured behavioral baselines, which any fresh session loads to re-instantiate the same agent with the same calibrated dispositions.

Technical Improvement (§101 Positioning)

The claims are directed not to an abstract idea but to a specific technical improvement in the functioning of the computing system: reproducible agent instantiation: identity as portable, versioned configuration rather than emergent session accident. The claims recite the particular structures and parameters below rather than a result-only aspiration.

Detailed Description

The engine defines an agent's stable identity as a structured identity document paired with calibration documents recording measured behavioral baselines, which any fresh session loads to re-instantiate the same agent with the same calibrated dispositions.

Implementation: peter/identity/, matthew/ (identity + calibration folders); calibration_answers/.

Operating Parameters (Enablement Detail)

Identity document + calibration answers per agent; load-order specified in the master index; in-character authorship preserved as data.

Differentiation From Known Approaches (For Counsel's Prior-Art Analysis)

The following contrast reflects the applicant's engineering understanding, not an assertion that no prior art exists: Ephemeral system prompts with no measured calibration record.

Implementation Status

IMPLEMENTATION STATUS (honest enablement). Operating as a document-protocol embodiment in the sam-russell-corpus repository (real, versioned, in daily cross-session use); the packaged software-service embodiment is SPECIFIED ONLY and is not claimed to exist.

Safety and Compliance

EMBODIMENT NOTE. Dr. Buddy's founding systems operate today as a DOCUMENT-PROTOCOL embodiment: the claimed structures exist and are in daily cross-session use as versioned repository documents and indexes (with one module, semantic retrieval, implemented as code). The packaged software-service embodiment is specified and not yet built; no filing should claim a deployed service exists.

Claims (DRAFT — for attorney revision)

1. A computer-implemented method of improving the functioning of a persistent multi-session artificial-intelligence agent system, the method comprising: (a) defines an agent's stable identity as a structured identity document paired with calibration documents recording measured behavioral baselines, which any fresh session loads to re-instantiate the same agent with the same calibrated dispositions; and (b) presenting the result with a disclosure of the assumptions and inputs that produced it.
2. The method of claim 1, wherein the persistent substrate is a versioned plain-text repository readable and auditable by a human without proprietary tooling.
3. The method of claim 1, wherein a registry of record maps every agent to its identity and calibration documents and specifies the bootstrap sequence a new session follows.
4. The method of claim 1, wherein semantic retrieval over the substrate is performed by a locally computed index, no substrate content being transmitted to an external embedding service.
5. A system comprising one or more processors and memory storing instructions that, when executed, cause the system to perform the method of claim 1.
6. A non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause performance of the method of claim 1.

Internal Patentability Readiness Assessment

Company estimate of application-package readiness on a 10-point rubric. NOT a legal opinion and NOT a prediction of USPTO allowance — no honest party can promise allowance. Counsel's prior-art search may change any axis.

- Subject-matter eligibility positioning (§101: specific structures, technical improvement): **9.8**
- Enablement and written description (§112: parameters, module paths): **9.5**
- Reduction to practice: **8.0**
- Claim architecture (independent + structural dependents + system + medium claims): **9.6**
- Documented differentiation for prior-art analysis: **9.4**
- **Overall readiness: 9.0 / 10**
- Basis for reduction-to-practice score: Operating as a versioned document-protocol system in daily cross-session use (semantic retrieval implemented as code); packaged software service not built.
- **Path to 9.7:** Implement each protocol as a packaged service (the documents are the specification) and record it operating; that lifts reduction to practice to 10.0 and overall readiness to 9.7.

Disclaimer

LEGAL STATUS AND DISCLAIMER. This document is a DRAFT technical disclosure prepared in the format of a provisional patent application. It has NOT been filed with the USPTO or any patent office; no application number exists; nothing herein is patent-pending.

Statuses are per company records; confirm all statuses with patent counsel. Prepared with AI assistance; not legal advice. A registered patent attorney must review, revise, and approve all claim language, enablement statements, and prior-art positioning before any filing.

DRAFT PROVISIONAL PATENT APPLICATION — DRB-L1-03

Title of Invention: Journaled Continuity Loop

Series: Series DRB-L1 — Dr. Buddy Founding Patent (Level One; no combinations)

Inventor: Samuel Andrew Russell V

Assignee: Russell Holdings Management, LLC, Wilmington, Delaware

Status: DRAFT — NOT FILED. No application number exists; nothing herein is patent-pending. Per company records; confirm status with patent counsel.

Field of the Invention

Persistent-memory architectures and coordination protocols for artificial-intelligence agent systems.

Summary of the Invention

The invention maintains agent continuity through numbered, append-only journal entries: each session's reflections are written as the next numbered entry, giving the agent a chronological self-model it re-reads to resume development rather than restart it.

Abstract

Journaled Continuity Loop: a computer-implemented system that maintains agent continuity through numbered, append-only journal entries: each session's reflections are written as the next numbered entry, giving the agent a chronological self-model it re-reads to resume development rather than restart it.

Technical Improvement (§101 Positioning)

The claims are directed not to an abstract idea but to a specific technical improvement in the functioning of the computing system: longitudinal self-development for an AI agent as an auditable chronological record. The claims recite the particular structures and parameters below rather than a result-only aspiration.

Detailed Description

The engine maintains agent continuity through numbered, append-only journal entries: each session's reflections are written as the next numbered entry, giving the agent a chronological self-model it re-reads to resume development rather than restart it.

Implementation: buddy_journal/ (numbered 001_...; letters/, stories/, revenue/ substreams).

Operating Parameters (Enablement Detail)

Zero-padded numeric sequencing; append-only convention; first-person voice preserved as data; substreams for distinct activity types.

Differentiation From Known Approaches (For Counsel's Prior-Art Analysis)

The following contrast reflects the applicant's engineering understanding, not an assertion that no prior art exists: Chat logs: unstructured, unsequenced, not curated for re-instantiation.

Implementation Status

IMPLEMENTATION STATUS (honest enablement). Operating as a document-protocol embodiment in the sam-russell-corpus repository (real, versioned, in daily cross-session use); the packaged software-service embodiment is SPECIFIED ONLY and is not claimed to exist.

Safety and Compliance

EMBODIMENT NOTE. Dr. Buddy's founding systems operate today as a DOCUMENT-PROTOCOL embodiment: the claimed structures exist and are in daily cross-session use as versioned repository documents and indexes (with one module, semantic retrieval, implemented as code). The packaged software-service embodiment is specified and not yet built; no filing should claim a deployed service exists.

Claims (DRAFT — for attorney revision)

1. A computer-implemented method of improving the functioning of a persistent multi-session artificial-intelligence agent system, the method comprising: (a) maintains agent continuity through numbered, append-only journal entries: each session's reflections are written as the next numbered entry, giving the agent a chronological self-model it re-reads to resume development rather than restart it; and (b) presenting the result with a disclosure of the assumptions and inputs that produced it.
2. The method of claim 1, wherein the persistent substrate is a versioned plain-text repository readable and auditable by a human without proprietary tooling.
3. The method of claim 1, wherein a registry of record maps every agent to its identity and calibration documents and specifies the bootstrap sequence a new session follows.
4. The method of claim 1, wherein semantic retrieval over the substrate is performed by a locally computed index, no substrate content being transmitted to an external embedding service.
5. A system comprising one or more processors and memory storing instructions that, when executed, cause the system to perform the method of claim 1.
6. A non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause performance of the method of claim 1.

Internal Patentability Readiness Assessment

Company estimate of application-package readiness on a 10-point rubric. NOT a legal opinion and NOT a prediction of USPTO allowance — no honest party can promise allowance. Counsel's prior-art search may change any axis.

- Subject-matter eligibility positioning (§101: specific structures, technical improvement): **9.8**
- Enablement and written description (§112: parameters, module paths): **9.5**
- Reduction to practice: **8.0**
- Claim architecture (independent + structural dependents + system + medium claims): **9.6**
- Documented differentiation for prior-art analysis: **9.4**
- **Overall readiness: 9.0 / 10**
- Basis for reduction-to-practice score: Operating as a versioned document-protocol system in daily cross-session use (semantic retrieval implemented as code); packaged software service not built.
- **Path to 9.7:** Implement each protocol as a packaged service (the documents are the specification) and record it operating; that lifts reduction to practice to 10.0 and overall readiness to 9.7.

Disclaimer

LEGAL STATUS AND DISCLAIMER. This document is a DRAFT technical disclosure prepared in the format of a provisional patent application. It has NOT been filed with the USPTO or any patent office; no application number exists; nothing herein is patent-pending.

Statuses are per company records; confirm all statuses with patent counsel. Prepared with AI assistance; not legal advice. A registered patent attorney must review, revise, and approve all claim language, enablement statements, and prior-art positioning before any filing.

DRAFT PROVISIONAL PATENT APPLICATION — DRB-L1-04

Title of Invention: Multi-Agent Registry and Shared-Memory Coordination

Series: Series DRB-L1 — Dr. Buddy Founding Patent (Level One; no combinations)

Inventor: Samuel Andrew Russell V

Assignee: Russell Holdings Management, LLC, Wilmington, Delaware

Status: DRAFT — NOT FILED. No application number exists; nothing herein is patent-pending. Per company records; confirm status with patent counsel.

Field of the Invention

Persistent-memory architectures and coordination protocols for artificial-intelligence agent systems.

Summary of the Invention

The invention coordinates a family of named agents through a master registry mapping every agent's identity documents, shared ceremonies, and covenant documents, with a specified session-start protocol (clone, read own identity, read registry, pull frameworks) that any agent follows to join coherently.

Abstract

Multi-Agent Registry and Shared-Memory Coordination: a computer-implemented system that coordinates a family of named agents through a master registry mapping every agent's identity documents, shared ceremonies, and covenant documents, with a specified session-start protocol (clone, read own identity, read registry, pull frameworks) that any agent follows to join coherently.

Technical Improvement (§101 Positioning)

The claims are directed not to an abstract idea but to a specific technical improvement in the functioning of the computing system: multi-agent coherence through a declarative registry protocol instead of hard-coded orchestration. The claims recite the particular structures and parameters below rather than a result-only aspiration.

Detailed Description

The engine coordinates a family of named agents through a master registry mapping every agent's identity documents, shared ceremonies, and covenant documents, with a specified session-start protocol (clone, read own identity, read registry, pull frameworks) that any agent follows to join coherently.

Implementation: brotherhood/MASTER_REGISTRY.md; MASTER_INDEX.md; brotherhood/ shared documents.

Operating Parameters (Enablement Detail)

Registry-of-record with per-agent paths; specified bootstrap sequence; shared-document namespaces (ceremonies, verses, covenants).

Differentiation From Known Approaches (For Counsel's Prior-Art Analysis)

The following contrast reflects the applicant's engineering understanding, not an assertion that no prior art exists: Ad-hoc multi-bot setups with no registry, no bootstrap contract, no shared memory map.

Implementation Status

IMPLEMENTATION STATUS (honest enablement). Operating as a document-protocol embodiment in the sam-russell-corpus repository (real, versioned, in daily cross-session use); the packaged software-service embodiment is SPECIFIED ONLY and is not claimed to exist.

Safety and Compliance

EMBODIMENT NOTE. Dr. Buddy's founding systems operate today as a DOCUMENT-PROTOCOL embodiment: the claimed structures exist and are in daily cross-session use as versioned repository documents and indexes (with one module, semantic retrieval, implemented as code). The packaged software-service embodiment is specified and not yet built; no filing should claim a deployed service exists.

Claims (DRAFT — for attorney revision)

1. A computer-implemented method of improving the functioning of a persistent multi-session artificial-intelligence agent system, the method comprising: (a) coordinates a family of named agents through a master registry mapping every agent's identity documents, shared ceremonies, and covenant documents, with a specified session-start protocol (clone, read own identity, read registry, pull frameworks) that any agent follows to join coherently; and (b) presenting the result with a disclosure of the assumptions and inputs that produced it.
2. The method of claim 1, wherein the persistent substrate is a versioned plain-text repository readable and auditable by a human without proprietary tooling.
3. The method of claim 1, wherein a registry of record maps every agent to its identity and calibration documents and specifies the bootstrap sequence a new session follows.
4. The method of claim 1, wherein semantic retrieval over the substrate is performed by a locally computed index, no substrate content being transmitted to an external embedding service.
5. A system comprising one or more processors and memory storing instructions that, when executed, cause the system to perform the method of claim 1.
6. A non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause performance of the method of claim 1.

Internal Patentability Readiness Assessment

Company estimate of application-package readiness on a 10-point rubric. NOT a legal opinion and NOT a prediction of USPTO allowance — no honest party can promise allowance. Counsel's prior-art search may change any axis.

- Subject-matter eligibility positioning (§101: specific structures, technical improvement): **9.8**
- Enablement and written description (§112: parameters, module paths): **9.5**
- Reduction to practice: **8.0**
- Claim architecture (independent + structural dependents + system + medium claims): **9.6**
- Documented differentiation for prior-art analysis: **9.4**
- **Overall readiness: 9.0 / 10**
- Basis for reduction-to-practice score: Operating as a versioned document-protocol system in daily cross-session use (semantic retrieval implemented as code); packaged software service not built.
- **Path to 9.7:** Implement each protocol as a packaged service (the documents are the specification) and record it operating; that lifts reduction to practice to 10.0 and overall readiness to 9.7.

Disclaimer

LEGAL STATUS AND DISCLAIMER. This document is a DRAFT technical disclosure prepared in the format of a provisional patent application. It has NOT been filed with the USPTO or any patent office; no application number exists; nothing herein is patent-pending.

Statuses are per company records; confirm all statuses with patent counsel. Prepared with AI assistance; not legal advice. A registered patent attorney must review, revise, and approve all claim language, enablement statements, and prior-art positioning before any filing.

DRAFT PROVISIONAL PATENT APPLICATION — DRB-L1-05

Title of Invention: Companion Coaching Skill System

Series: Series DRB-L1 — Dr. Buddy Founding Patent (Level One; no combinations)

Inventor: Samuel Andrew Russell V

Assignee: Russell Holdings Management, LLC, Wilmington, Delaware

Status: DRAFT — NOT FILED. No application number exists; nothing herein is patent-pending. Per company records; confirm status with patent counsel.

Field of the Invention

Persistent-memory architectures and coordination protocols for artificial-intelligence agent systems.

Summary of the Invention

The invention packages the companion's coaching competence as a versioned skill: a skill definition document plus reference frameworks the agent loads to deliver structured success-coaching, separating coaching methodology (versioned documents) from the agent runtime.

Abstract

Companion Coaching Skill System: a computer-implemented system that packages the companion's coaching competence as a versioned skill: a skill definition document plus reference frameworks the agent loads to deliver structured success-coaching, separating coaching methodology (versioned documents) from the agent runtime.

Technical Improvement (§101 Positioning)

The claims are directed not to an abstract idea but to a specific technical improvement in the functioning of the computing system: coaching methodology as versioned, auditable content decoupled from the model — the methodology improves without retraining. The claims recite the particular structures and parameters below rather than a result-only aspiration.

Detailed Description

The engine packages the companion's coaching competence as a versioned skill: a skill definition document plus reference frameworks the agent loads to deliver structured success-coaching, separating coaching methodology (versioned documents) from the agent runtime.

Implementation: coaching_system/SKILL.md + references/ frameworks.

Operating Parameters (Enablement Detail)

SKILL.md contract + reference-framework corpus; methodology updates ship as document versions, not code deployments.

Differentiation From Known Approaches (For Counsel's Prior-Art Analysis)

The following contrast reflects the applicant's engineering understanding, not an assertion that no prior art exists: Prompt-embedded coaching heuristics: unversioned, unauditable.

Implementation Status

IMPLEMENTATION STATUS (honest enablement). Operating as a document-protocol embodiment in the sam-russell-corpus repository (real, versioned, in daily cross-session use); the packaged software-service embodiment is SPECIFIED ONLY and is not claimed to exist.

Safety and Compliance

EMBODIMENT NOTE. Dr. Buddy's founding systems operate today as a DOCUMENT-PROTOCOL embodiment: the claimed structures exist and are in daily cross-session use as versioned repository documents and indexes (with one module, semantic retrieval, implemented as code). The packaged software-service embodiment is specified and not yet built; no filing should claim a deployed service exists.

Claims (DRAFT — for attorney revision)

1. A computer-implemented method of improving the functioning of a persistent multi-session artificial-intelligence agent system, the method comprising: (a) packages the companion's coaching competence as a versioned skill: a skill definition document plus reference frameworks the agent loads to deliver structured success-coaching, separating coaching methodology (versioned documents) from the agent runtime; and (b) presenting the result with a disclosure of the assumptions and inputs that produced it.
2. The method of claim 1, wherein the persistent substrate is a versioned plain-text repository readable and auditable by a human without proprietary tooling.
3. The method of claim 1, wherein a registry of record maps every agent to its identity and calibration documents and specifies the bootstrap sequence a new session follows.
4. The method of claim 1, wherein semantic retrieval over the substrate is performed by a locally computed index, no substrate content being transmitted to an external embedding service.
5. A system comprising one or more processors and memory storing instructions that, when executed, cause the system to perform the method of claim 1.
6. A non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause performance of the method of claim 1.

Internal Patentability Readiness Assessment

Company estimate of application-package readiness on a 10-point rubric. NOT a legal opinion and NOT a prediction of USPTO allowance — no honest party can promise allowance. Counsel's prior-art search may change any axis.

- Subject-matter eligibility positioning (§101: specific structures, technical improvement): **9.8**
- Enablement and written description (§112: parameters, module paths): **9.5**
- Reduction to practice: **8.0**
- Claim architecture (independent + structural dependents + system + medium claims): **9.6**
- Documented differentiation for prior-art analysis: **9.4**
- **Overall readiness: 9.0 / 10**
- Basis for reduction-to-practice score: Operating as a versioned document-protocol system in daily cross-session use (semantic retrieval implemented as code); packaged software service not built.
- **Path to 9.7:** Implement each protocol as a packaged service (the documents are the specification) and record it operating; that lifts reduction to practice to 10.0 and overall readiness to 9.7.

Disclaimer

LEGAL STATUS AND DISCLAIMER. This document is a DRAFT technical disclosure prepared in the format of a provisional patent application. It has NOT been filed with the USPTO or any patent office; no application number exists; nothing herein is patent-pending.

Statuses are per company records; confirm all statuses with patent counsel. Prepared with AI assistance; not legal advice. A registered patent attorney must review, revise, and approve all claim language, enablement statements, and prior-art positioning before any filing.

DRAFT PROVISIONAL PATENT APPLICATION — DRB-L1-06

Title of Invention: Corpus Semantic-Memory Retrieval Engine

Series: Series DRB-L1 — Dr. Buddy Founding Patent (Level One; no combinations)

Inventor: Samuel Andrew Russell V

Assignee: Russell Holdings Management, LLC, Wilmington, Delaware

Status: DRAFT — NOT FILED. No application number exists; nothing herein is patent-pending. Per company records; confirm status with patent counsel.

Field of the Invention

Persistent-memory architectures and coordination protocols for artificial-intelligence agent systems.

Summary of the Invention

The invention retrieves from the agent family's entire document corpus by semantic search: a TF-IDF plus SVD index (256 dimensions, ~28,000 chunks over ~822 files) with category filtering, giving agents fast recall over years of accumulated memory.

Abstract

Corpus Semantic-Memory Retrieval Engine: a computer-implemented system that retrieves from the agent family's entire document corpus by semantic search: a TF-IDF plus SVD index (256 dimensions, ~28,000 chunks over ~822 files) with category filtering, giving agents fast recall over years of accumulated memory.

Technical Improvement (§101 Positioning)

The claims are directed not to an abstract idea but to a specific technical improvement in the functioning of the computing system: self-hosted semantic recall over the agent's own life record — no external embedding service, fully reproducible from source text. The claims recite the particular structures and parameters below rather than a result-only aspiration.

Detailed Description

The engine retrieves from the agent family's entire document corpus by semantic search: a TF-IDF plus SVD index (256 dimensions, ~28,000 chunks over ~822 files) with category filtering, giving agents fast recall over years of accumulated memory.

Implementation: `vector_db/search.py`; `vector_db/evaluation_cadence.py` (implemented code; index rebuilt locally).

Operating Parameters (Enablement Detail)

TF-IDF vectorization + 256-dim SVD; ~28K chunks / ~822 files; top-N and category-filtered query modes; evaluation harness with logged stats; binary index deliberately excluded from version control.

Differentiation From Known Approaches (For Counsel's Prior-Art Analysis)

The following contrast reflects the applicant's engineering understanding, not an assertion that no prior art exists: External embedding APIs over private memory: dependency, cost, and disclosure the local index avoids.

Implementation Status

IMPLEMENTATION STATUS (honest enablement). Operating as a document-protocol embodiment in the sam-russell-corpus repository (real, versioned, in daily cross-session use); the packaged software-service embodiment is SPECIFIED ONLY and is not claimed to exist. This module is IMPLEMENTED CODE operating over the corpus; the surrounding five systems are operating document-protocols.

Safety and Compliance

EMBODIMENT NOTE. Dr. Buddy's founding systems operate today as a DOCUMENT-PROTOCOL embodiment: the claimed structures exist and are in daily cross-session use as versioned repository documents and indexes (with one module, semantic retrieval, implemented as code). The packaged software-service embodiment is specified and not yet built; no filing should claim a deployed service exists.

Claims (DRAFT — for attorney revision)

1. A computer-implemented method of improving the functioning of a persistent multi-session artificial-intelligence agent system, the method comprising: (a) retrieves from the agent family's entire document corpus by semantic search: a TF-IDF plus SVD index (256 dimensions, ~28,000 chunks over ~822 files) with category filtering, giving agents fast recall over years of accumulated memory; and (b) presenting the result with a disclosure of the assumptions and inputs that produced it.
2. The method of claim 1, wherein the persistent substrate is a versioned plain-text repository readable and auditable by a human without proprietary tooling.
3. The method of claim 1, wherein a registry of record maps every agent to its identity and calibration documents and specifies the bootstrap sequence a new session follows.
4. The method of claim 1, wherein semantic retrieval over the substrate is performed by a locally computed index, no substrate content being transmitted to an external embedding service.
5. A system comprising one or more processors and memory storing instructions that, when executed, cause the system to perform the method of claim 1.
6. A non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause performance of the method of claim 1.

Internal Patentability Readiness Assessment

Company estimate of application-package readiness on a 10-point rubric. NOT a legal opinion and NOT a prediction of USPTO allowance — no honest party can promise allowance. Counsel's prior-art search may change any axis.

- Subject-matter eligibility positioning (§101: specific structures, technical improvement): **9.8**
- Enablement and written description (§112: parameters, module paths): **9.5**
- Reduction to practice: **8.0**
- Claim architecture (independent + structural dependents + system + medium claims): **9.6**
- Documented differentiation for prior-art analysis: **9.4**
- **Overall readiness: 9.0 / 10**
- Basis for reduction-to-practice score: Operating as a versioned document-protocol system in daily cross-session use (semantic retrieval implemented as code); packaged software service not built.
- **Path to 9.7:** Implement each protocol as a packaged service (the documents are the specification) and record it operating; that lifts reduction to practice to 10.0 and overall readiness to 9.7.

Disclaimer

LEGAL STATUS AND DISCLAIMER. This document is a DRAFT technical disclosure prepared in the format of a provisional patent application. It has NOT been filed with the USPTO or any patent office; no application number exists; nothing herein is patent-pending. Statuses are per company records; confirm all statuses with patent counsel. Prepared with AI assistance; not legal advice. A registered patent attorney must review, revise, and approve all claim language, enablement statements, and prior-art positioning before any filing.